

# Optimization problems constrained by parameter sums

John C. Nash, retired professor, Telfer School of Management, University of Ottawa

Gabor Grothendieck, GKX Group

Ravi Varadhan, Johns Hopkins University Medical School

2023-7-31

## Abstract

This article presents a discussion of optimization problems where the objective function  $f(\mathbf{x})$  has parameters that are constrained by some scaling, so that  $q(\mathbf{x}) = \text{constant}$ , where this function  $q()$  involves a sum of the parameters, their squares, or similar simple function. Our focus is on ways to use standardized optimization programs to solve such problems rather than specialized codes.

## Background

We consider problems where we want to minimize or maximize a function subject to a constraint that the sum of some function of the parameters, e.g., their sum of squares, must equal some constant. Because these problems all have an objective that is dependent on a scaled set of parameters where the scale is defined by a sum, sum of squares, or similar sum of the parameters, we will refer to them as **sumscale** optimization problems.

We have observed questions about problems like this on the R-help mailing list:

Jul 19, 2012 at 10:24 AM, Linh Tran <Tranlm@berkeley.edu> wrote:

```
> Hi fellow R users,
>
> I am desperately hoping there is an easy way to do this in R.
>
> Say I have three functions:
>
> f(x) = x^2
> f(y) = 2y^2
> f(z) = 3z^2
>
> constrained such that x+y+z=c (let c=1 for simplicity).
>
> I want to find the values of x,y,z that will minimize
f(x) + f(y) + f(z).
```

If the parameters  $x$ ,  $y$  and  $z$  are non-negative, this problem can actually be solved as a Quadratic Program. We revisit this problem at the end of this article.

Other examples of this type of objective function are:

- The maximum volume of a regular polyhedron where the sum of the lengths of the sides is fixed.
- The minimum negative log likelihood for a multinomial model.
- The Rayleigh Quotient for the maximal or minimal eigensolutions of a matrix, where the eigenvectors should be normalized so the square norm of the vector is 1.

- The minimum of the extended Rosenbrock function of the form given by the `adagio` package (Borchers (2022)) on the unit ball, that is, where the sum of squares of the parameters is 1.

For the moment, let us consider a basic example, which is

**Problem A: Minimize  $(-\prod \mathbf{x})$  subject to  $\sum \mathbf{x} = 1$**

This is a very simplified version of the multinomial maximum likelihood problem.

## Difficulties using general optimization with sumscales problems

Let us use the basic example above to consider how we might formulate Problem A for a computational solution in R.

One possibility is to select one of the parameters and solve for it in terms of the others. Let this be the last parameter  $x_n$ , so that the set of parameters to be optimized is  $\mathbf{y} = (x_1, x_1, \dots, x_{n-1})$  where  $n$  is the original size of our problem. We now have the unconstrained problem

$$\text{minimize}(-(\prod \mathbf{y}) * (1 - \sum \mathbf{y}))$$

This is easily coded and tried. We will use a very simple start, namely, the sequence  $1, 2, \dots, (n-1)$  scaled by  $1/n^2$ . We will also specify that the gradient is to be computed by a central approximation (see `gcentral.R` from package `optimx`).

```
require(optimx, quietly=TRUE)
pr <- function(y) {
  - prod(y)*(1-sum(y))
}
cat("test the simple product for n=5\n")

## test the simple product for n=5

meth <- c("Nelder-Mead", "BFGS")
n<-5
st<-1:(n-1)/(n*n)
ans<-opm(st, pr, gr="gcentral", control=list(trace=0))
ao<-summary(ans, order=value)
print(ao)
```

##	p1	s1	p2	s2	p3	s3	p4	s4	value	fevals
## BFGS	0.2000000	0.1999986	0.2000037	0.1999985	-0.00032	102				
## Nelder-Mead	0.2000034	0.1999983	0.2000017	0.2000021	-0.00032	331				
##	gevals	hevals	conv	kkt1	kkt2	xtime				
## BFGS	96	0	0	TRUE	TRUE	0.002				
## Nelder-Mead	0	0	0	TRUE	TRUE	0.016				

While these codes work fine for small  $n$ , it is fairly easy to see that there are computational problems as the size of the problem increases. Since the sum of the parameters is constrained to be equal to 1, the parameters are of the order of  $1/n$ , and the function therefore of the order of  $1/(n^n)$ , which underflows around  $n = 144$  in R.

## Other formulations

Traditionally, statisticians solve maximum likelihood problems by **minimizing** the negative log-likelihood. That is, the objective function is formed as (-1) times the logarithm of the likelihood. This converts our product to a sum. Choosing the first parameter to be the one determined by the summation constraint, we can write the function and gradient quite easily. As programs that try to find the minimum may change the parameters so that logarithms of non-positive numbers are attempted, we have put some safeguards in the

function `nll`. At this point we have assumed the gradient calculation is only attempted if the function can be computed satisfactorily, so we have not put safeguards in the gradient.

```
nll <- function(y) {
  if ((any(y <= 10*.Machine$double.xmin)) || (sum(y)>1-.Machine$double.eps))
    .Machine$double.xmax
  else - sum(log(y)) - log(1-sum(y))
}
nll.g <- function(y) { - 1/y + 1/(1-sum(y))} # so far not safeguarded
n<-5
x0<-(2:n)/n^2
library(numDeriv)
dx0<-nll.g(x0)
dx0n<-grad(nll, x0)
cat("Max Abs diff between analytic and approx. gradient =",max(abs(dx0-dx0n)), "\n")
```

```
## Max Abs diff between analytic and approx. gradient = 5.92224e-10
```

We can easily try several optimization methods using the `optimx` package. Here are the calls, which overall did not perform as well as we would like. Note that we do not ask for `method="ALL"` as we found that some of the methods, in particular those using Powell's quadratic approximation methods, seem to get "stuck". Instead, we have specified a list of methods `mset`, though some of these also run into scaling problems.

```
require(optimx, quietly=TRUE)
mset<-c("L-BFGS-B", "BFGS", "CG", "spg", "ucminf", "nlm", "nlminb", "nvm", "ncg", "tnewt")
# numerical approximation using forward difference
a5<-opm(2:n/n^2, nll, gr="grfwd", method=mset,
  control=list(dowarn=FALSE, kkt=FALSE, trace=0))
```

```
## Error in optim(par = par, fn = efn, gr = egr, method = method, hessian = FALSE, :
## non-finite value supplied by optim
## Gradient check details: max. relative difference in gradients= 1.100932
##
## analytic gradient: -10.22727 -12.87881 -13.63639 -13.6364
##
## numerical gradient: -10.22727 -6.060606 -3.977273 -2.727273Error in BB::spg(par = spar, fn = efn, g
## Analytic gradient does not seem correct! See comparison above. Fix it, remove it, or increase check
## Error in nlm(f = fghfn, p = spar, iterlim = iterlim, print.level = print.level) :
## probable coding error in analytic gradient
```

```
summary(a5, order=value)
```

	p1	s1	p2	s2	p3	s3	p4	s4	value
## nvm	0.2000002		0.2000000		0.1999998		0.1999996		8.047190e+00
## tnewt	0.1999999		0.1999999		0.2000002		0.1999995		8.047190e+00
## ncg	0.2000008		0.2000006		0.1999990		0.1999996		8.047190e+00
## CG	0.2000007		0.2000008		0.1999991		0.1999989		8.047190e+00
## nlminb	0.2000054		0.1999971		0.1999976		0.1999975		8.047190e+00
## BFGS	0.1999961		0.1999887		0.2000367		0.2000092		8.047190e+00
## ucminf	0.1736718		0.1929237		0.2013754		0.2203217		8.063850e+00
## L-BFGS-B	NA		NA		NA		NA		8.988466e+307
## spg	NA		NA		NA		NA		8.988466e+307
## nlm	NA		NA		NA		NA		8.988466e+307
##	fevals	gevals	hevals	conv	kkt1	kkt2	xtime		
## nvm	130	36	0	0	NA	NA	0.005		
## tnewt	167	166	0	0	NA	NA	0.007		
## ncg	2245	314	0	1	NA	NA	0.017		

```
## CG          362    94      0    0    NA    NA 0.003
## nlminb       43    29      0    0    NA    NA 0.001
## BFGS        231    31      0    0    NA    NA 0.001
## ucminf       12    12      0    0    NA    NA 0.003
## L-BFGS-B      2     2      0 9999    NA    NA 0.000
## spg          35     1      0 9999    NA    NA 0.004
## nlm          6     6      0 9999    NA    NA 0.000
```

```
# analytical gradient
a5g<-opm(2:n/n^2, nll, nll.g, method=mset,
         control=list(dowarn=FALSE,kkt=FALSE, trace=0))
```

```
## Error in optim(par = par, fn = efn, gr = egr, method = method, hessian = FALSE, :
## non-finite value supplied by optim
```

```
summary(a5g, order=value)
```

##	p1 s1	p2 s2	p3 s3	p4 s4	value
## ucminf	0.2000000	0.2000000	0.2000000	0.2000000	8.047190e+00
## nvm	0.2000000	0.2000000	0.2000000	0.2000000	8.047190e+00
## ncg	0.2000000	0.2000000	0.2000000	0.2000000	8.047190e+00
## tnewt	0.2000000	0.2000000	0.2000000	0.2000000	8.047190e+00
## spg	0.2000000	0.2000000	0.2000000	0.2000000	8.047190e+00
## CG	0.2000000	0.2000000	0.2000000	0.2000000	8.047190e+00
## nlm	0.2000006	0.1999995	0.2000000	0.2000000	8.047190e+00
## BFGS	0.2000007	0.1999989	0.2000012	0.1999981	8.047190e+00
## nlminb	0.2000004	0.1999990	0.1999989	0.1999992	8.047190e+00
## L-BFGS-B	NA	NA	NA	NA	8.988466e+307

  

##	fevals	gevals	hevals	conv	kkt1	kkt2	xtime
## ucminf	14	14	0	0	NA	NA	0.001
## nvm	31	12	0	0	NA	NA	0.001
## ncg	29	12	0	0	NA	NA	0.001
## tnewt	27	26	0	0	NA	NA	0.001
## spg	51	15	0	0	NA	NA	0.001
## CG	59	21	0	0	NA	NA	0.000
## nlm	19	19	0	0	NA	NA	0.001
## BFGS	33	9	0	0	NA	NA	0.000
## nlminb	24	12	0	0	NA	NA	0.000
## L-BFGS-B	2	2	0	9999	NA	NA	0.003

```
# analytical gradient and bounds on parameters
a5gb<-opm(2:n/n^2, nll, nll.g, lower=0, upper=1, method=mset,
          control=list(dowarn=FALSE,kkt=FALSE,trace=0))
```

```
## Warning in opm(2:n/n^2, nll, nll.g, lower = 0, upper = 1, method = mset, :
## method requested does not handle bounds
```

```
## Error in optim(par = par, fn = efn, gr = egr, lower = lower, upper = upper, :
## non-finite value supplied by optim
```

```
summary(a5gb, order=value)
```

##	p1 s1	p2 s2	p3 s3	p4 s4	value
## nvm	0.2000000	0.2000000	0.2000000	0.2000000	8.047190e+00
## tnewt	0.2000000	0.2000000	0.2000000	0.2000000	8.047190e+00
## ncg	0.2000000	0.2000000	0.2000000	0.2000000	8.047190e+00
## nlminb	0.2000004	0.1999999	0.1999989	0.1999992	8.047190e+00
## L-BFGS-B	NA	NA	NA	NA	8.988466e+307

##	fevals	gevals	hevals	conv	kkt1	kkt2	xtime
## nvm	28	14	0	0	NA	NA	0.002
## tnewt	20	19	0	0	NA	NA	0.001
## ncg	26	12	0	0	NA	NA	0.000
## nlminb	24	12	0	0	NA	NA	0.000
## L-BFGS-B	2	2	0	9999	NA	NA	0.000

Most, but not all, of the methods find the solution for the  $n = 5$  case. The exception (L-BFGS-B) is due to the optimization method trying to compute the gradient where  $\text{sum}(\mathbf{x})$  is greater than 1. We have not tried to determine the source of this particular issue. However, it is almost certainly a consequence of too large a step. The particular form of  $\log(1 - \text{sum}(\mathbf{x}))$  is undefined once the argument of the logarithm is negative. Indeed, this is the basis of logarithmic barrier functions for constraints. There is a similar issue with the  $n - 1$  parameters near zero. Negative values will cause difficulties.

Numerical gradient approximations will similarly fail, particularly as step sizes are often of the order of  $1\text{E-}7$  in size. There is generally no special check within numerical gradient routines to apply bounds. Note also that a lower bound of 0 on parameters is not adequate, since  $\log(0)$  is undefined. Choosing a bound large enough to avoid the logarithm of a zero or negative argument while still being small enough to allow for parameter optimization is non-trivial.

## Transformed problems or parameters

When problems give difficulties, it is common to re-formulate them by transformations of the function or the parameters.

### Using a projection

Objective functions defined by  $(-1) * \prod \mathbf{x}$  or  $(-1) * \sum \log(\mathbf{x})$  will change with the scale of the parameters. Moreover, the constraint  $\sum \mathbf{x} = 1$  effectively imposes the scaling

$$\mathbf{x}_{scaled} = \mathbf{x} / \sum \mathbf{x}$$

The optimizer `spg` from package `BB` allows us to project our search direction to satisfy constraints. Thus, we could use the following approach.

```
require(BB, quietly=TRUE)
nllrv <- function(x) {- sum(log(x))}
nllrv.g <- function(x) {- 1/x }
proj <- function(x) {x/sum(x)}
n <- 5
aspg <- spg(par=(1:n)/n^2, fn=nllrv, gr=nllrv.g, project=proj,
            control=list(trace=TRUE, triter=1))
```

```
## iter: 0 f-value: 11.30689 pgrad: 0.3607565
## iter: 1 f-value: 8.881109 pgrad: 0.3554228
## iter: 2 f-value: 8.095091 pgrad: 0.09254548
## iter: 3 f-value: 8.048691 pgrad: 0.018282
## iter: 4 f-value: 8.047191 pgrad: 0.0004657523
## iter: 5 f-value: 8.04719 pgrad: 7.331566e-06
```

```
aspgn <- spg(par=(1:n)/n^2, fn=nllrv, project=proj,
             control=list(trace=TRUE, triter=1)) # using internal grad approx.
```

```
## iter: 0 f-value: 11.30689 pgrad: 0.1333334
## iter: 1 f-value: 8.04719 pgrad: 1.225824e-07
```

```
cat("F_optimal: with gradient=",aspg$value," num. approx.=",aspgn$value,"\n")
```

```
## F_optimal: with gradient= 8.04719 num. approx.= 8.04719
```

```
pbest<-rep(1/n, n)
```

```
cat("fbest = ",nllrv(pbest)," when all parameters = ", pbest[1],"\n")
```

```
## fbest = 8.04719 when all parameters = 0.2
```

```
cat("deviations: with gradient=",max(abs(aspg$par-pbest)),  
    " num. approx.=",max(abs(aspgn$par-pbest)), "\n")
```

```
## deviations: with gradient= 3.81244e-06 num. approx.= 3.81244e-06
```

Here the projection `proj` is the key to success of method `spg`. Other methods (as yet) do not have the flexibility to impose the projection directly. We would need to carefully build the projection into the function(s) and/or the method codes. This was done by Geradin (1971) for the Rayleigh quotient problem, but requires a number of changes to the program code. Why `spg()` does (for this case) much better using the internal numerical gradient approximation than the analytic gradient is an open question.

### `log()` transformation of parameters

A common method to ensure parameters are positive is to transform them. In the present case, optimizing over parameters that are the logarithms of the parameters above ensures we have positive arguments to most of the elements of the negative log likelihood. Here is the code. Note that the parameters used in optimization are “`lx`” and not `x`.

```
enll <- function(lx) {  
  x<-exp(lx)  
  fval<- - sum( log( x/sum(x) ) )  
}  
enll.g <- function(lx){  
  x<-exp(lx)  
  g<-length(x)/sum(x) - 1/x  
  gval<-g*exp(lx)  
}
```

But where is our constraint? Here we have noted that we could define the objective function only to within the scaling  $x/\sum(x)$ . There is a minor nuisance, in that we need to re-scale our parameters after solution to have them in a standard form. This is most noticeable if one uses `optimx` and displays the results of `all.methods`. In the following, we extract the best solution for the 5-parameter problem.

```
require(optimx, quietly=TRUE) # just to be sure  
st<-1:5/10 # 5 parameters, crude scaling to start  
a5x<-opm(st, enll, enll.g, method="MOST", control=list(trace=0))  
a5x<-a5x[order(a5x$value),]  
cat("Proposed best solution has minimum=",a5x[1,length(st)+1],"\n")
```

```
## Proposed best solution has minimum= 8.04719
```

```
cat("Coeffs:")
```

```
## Coeffs:
```

```
print(a5x[1:length(st)])
```

```
##          p1 p2 p3 p4 p5  
## nlminb 0.3 0.3 0.3 0.3 0.3
```

While there are reasons to think that the indeterminacy might upset the optimization codes, in practice, the objective and gradient above are generally well-behaved, though they did reveal that tests of the size of the gradient used, in particular, to decide to terminate iterations in `Rcgmin()` were too hasty in stopping progress for problems with larger numbers of parameters. A user-specified tolerance is now allowed; for example `control=list(tol=1e-12)`.

Let us try a larger problem in 100 parameters.

```
require(optimx, quietly=TRUE)
st<-1:100/1e3 # large
stenll<-enll(st)
cat("Initial function value =",stenll,"\n")

## Initial function value = 460.5587

tym<-system.time(acgbig<-Rcgmin(st, enll, enll.g,
                                control=list(trace=0, tol=1e-32)))[[3]]
cat("Time = ",tym," fval=",acgbig$value,"\n")

## Time = 0.017 fval= 460.517

xnor<-acgbig$par/sum(acgbig$par)
print(xnor)

## [1] 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01
## [16] 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01
## [31] 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01
## [46] 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01
## [61] 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01
## [76] 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01
## [91] 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01
```

One worrying aspect of the solution is that the objective function at the start and end differ by a tiny amount.

### Another transformation

A slightly different transformation or projection is inspired by spherical coordinates.

```
library(optimx)
proj2 <- function(theta) {
  theta2 <- theta^2
  s2 <- theta2 / (1 + theta2)
  cumprod(c(1, s2)) * c(1-s2, 1)
}
obj <- function(theta) - sum(log(proj2(theta)))
n <- 5
ans <- spg(seq(n-1), obj)

## iter: 0 f-value: 11.15175 pgrad: 3
## iter: 10 f-value: 8.78015 pgrad: 0.5806909
## iter: 20 f-value: 8.04719 pgrad: 3.925749e-06

proj2(ans$par)

## [1] 0.2000000 0.2000007 0.2000002 0.1999996 0.1999995

n<-100
# check
obj(seq(n-1))
```

```
## [1] 843.9598
```

```
ans100 <- spg(seq(n-1), obj, control=list(trace=FALSE), quiet=TRUE)
ans100$value
```

```
## [1] 460.517
```

```
proj2( (ans100$par) )
```

```
## [1] 0.010000000 0.010000002 0.009999998 0.009999998 0.010000001 0.010000001
## [7] 0.009999999 0.010000001 0.010000001 0.010000001 0.010000001 0.009999999
## [13] 0.010000002 0.010000001 0.010000001 0.009999999 0.009999999 0.010000002
## [19] 0.010000001 0.010000000 0.010000003 0.010000002 0.010000003 0.010000003
## [25] 0.010000002 0.010000000 0.010000001 0.010000003 0.010000000 0.010000000
## [31] 0.010000003 0.009999998 0.010000001 0.010000002 0.009999999 0.009999999
## [37] 0.009999998 0.010000000 0.009999999 0.009999998 0.010000001 0.010000002
## [43] 0.009999999 0.010000001 0.010000002 0.010000000 0.010000002 0.009999999
## [49] 0.010000004 0.010000002 0.010000002 0.009999999 0.009999998 0.010000001
## [55] 0.010000002 0.010000002 0.009999999 0.010000000 0.010000000 0.010000000
## [61] 0.010000002 0.010000000 0.009999999 0.009999999 0.010000000 0.009999999
## [67] 0.009999999 0.009999999 0.009999999 0.009999999 0.009999999 0.009999999
## [73] 0.009999999 0.009999999 0.009999999 0.009999999 0.009999999 0.009999999
## [79] 0.010000000 0.009999999 0.009999999 0.009999999 0.009999998 0.010000000
## [85] 0.009999999 0.009999999 0.009999999 0.009999999 0.010000000 0.009999999
## [91] 0.009999999 0.010000000 0.009999999 0.009999988 0.010000001 0.010000001
## [97] 0.010000001 0.010000007 0.010000018 0.009999973
```

Since this transformation is embedded into the objective function, we could run all the optimizers in `optimx` as follows. This takes some time, as the derivative-free methods appear to have more difficulty with this formulation. Moreover, `Rcgmin` and `Rvmmmin` are not recommended when an analytic gradient is not provided.

```
# turn off kkt to save time
tmeth<-c("nCG", "nvm", "lbfgs", "ucminf", "bobyqa", "tnewt", "slsqp")
mfv<-10*(n-1)^2 # set fn eval count big enough to avoid commonArgs error (bobyqa)
allansf<- opm(seq(n-1)/n, obj, gr="grfwd", method=tmeth,
              control=list(kkt=FALSE, dowarn=FALSE,maxfeval=mfv))
summary(allansf, order = "list(round(value, 3), fevals)", par.select = FALSE)
```

	value	fevals	gevals	hevals	conv	kkt1	kkt2	xtime
## bobyqa	460.5170	8818	0	0	0	NA	NA	1.958
## lbfgs	494.4718	35	35	0	-1001	NA	NA	0.115
## slsqp	520.9416	2255	2254	0	0	NA	NA	1.285
## nvm	584.0714	67	6	0	3	NA	NA	0.006
## nCG	1784.4344	3	1	0	0	NA	NA	0.001
## tnewt	1818.2848	38	37	0	9999	NA	NA	0.020
## ucminf	13140.9297	9	9	0	0	NA	NA	0.011

```
allansc<- opm(seq(n-1)/n, obj, gr="grcentral", method=tmeth,
              control=list(kkt=FALSE, dowarn=FALSE,maxfeval=mfv))
summary(allansc, order = "list(round(value, 3), fevals)", par.select = FALSE)
```

	value	fevals	gevals	hevals	conv	kkt1	kkt2	xtime
## lbfgs	460.5170	42	42	0	0	NA	NA	0.176
## slsqp	460.5170	108	107	0	0	NA	NA	0.207
## ucminf	460.5170	773	773	0	0	NA	NA	1.640
## nvm	460.5170	1464	1356	0	0	NA	NA	7.385
## bobyqa	460.5170	8818	0	0	0	NA	NA	2.009
## tnewt	529.5788	4027	4026	0	0	NA	NA	4.230



```
## ncg      906.3752      3      1      0      0      NA      NA 0.001
allansp<- opm(seq(n-1)/n, obj, gr="grpracma", method=tmeth,
              control=list(kkt=FALSE, dowarn=FALSE,maxfeval=mfv))
summary(allansp, order = "list(round(value, 3), fevals)", par.select = FALSE)

##          value fevals gevals hevals conv kkt1 kkt2 xtime
## lbfgs  460.5170     42     42      0      0      NA      NA 0.126
## slsqp  460.5170    108    107      0      0      NA      NA 0.217
## ucminf 460.5170    756    756      0      0      NA      NA 1.616
## nvm    460.5170   1663   1550      0      0      NA      NA 8.920
## bobyqa 460.5170   8818      0      0      0      NA      NA 2.012
## tnewt  548.8028   4079   4078      0      0      NA      NA 4.550
## ncg    906.3752      3      1      0      0      NA      NA 0.002
```

Note that the more accurate gradient approximations appear to give slightly better results.

### Use the gradient equations

Another approach is to “solve” the gradient equations. We can do this with a sum of squares minimizer, though the `nls` function in R is specifically NOT useful as it cannot, by default, deal with small or zero residuals. However, `nlf` from package `nlsr` is capable of dealing with such problems. Unfortunately, it will be slow as it has to generate the Jacobian by numerical approximation unless we can provide a function to prepare the Jacobian analytically. Moreover, the determination of the Jacobian is still subject to the unfortunate scaling issues we have been confronting throughout this article.

This approach is yet to be tried.

## The Rayleigh Quotient

The maximal and minimal eigensolutions of a symmetric matrix  $A$  are extrema of the Rayleigh Quotient

$$R(x) = (x'Ax)/(x'x)$$

We can also deal with generalized eigenproblems of the form

$$Ax = eBx$$

where  $B$  is symmetric and positive definite by using the Rayleigh Quotient

$$R_g(x) = (x'Ax)/(x'Bx)$$

Once again, the objective is scaled by the parameters, this time by their sum of squares. Alternatively, we may think of requiring the **normalized** eigensolution, which is given as

$$x_{normalized} = x/\sqrt{x'x}$$

We will first try the projected gradient method `spg` from `BB`. Below is the code, where our test uses a matrix called the Moler matrix Nash (1979, Appendix 1). This matrix is simple to generate and is positive definite, but has one small eigenvalue that may not be computed to high relative precision. That is, we may get a number that is precise relative to the largest eigenvalue, but having few trustworthy digits.

Let us set up some infrastructure for the Rayleigh Quotient of the matrix.

```

molerbuild<-function(n){ # Create the moler matrix of order n
  # A[i,j] = i for i=j, min(i,j)-2 otherwise
  A <- matrix(0, nrow = n, ncol = n)
  j <- 1:n
  for (i in 1:n) {
    A[i, 1:i] <- pmin(i, 1:i) - 2
  }
  A <- A + t(A)
  diag(A) <- 1:n
  A
}

raynum<-function(x, A){
  raynum<-as.numeric((t(x)%*%A)%*%x)
}

RQ <- function(x, A){
  RQ <- raynum(x, A)/sum(x*x)
}

# proj<-function(x) { x/sqrt(sum(x*x)) } # Original
proj <- function(x) {sign(x[1]) * x/sqrt(c(crossprod(x))) } # from ravi

# proj1 <- function(x) {sign(x[1]) * x/max(x)} # no use!

```

Let us now try to use `BB::spg()` to find the largest eigenvalue by minimizing the Rayleigh Quotient of  $-A$ .

```

require(BB, quietly=TRUE)
n<-10
set.seed(4321)
x<-runif(n)
AA<-molerbuild(n)
cat("Eigenvalues from eigen:")

## Eigenvalues from eigen:
meig<-eigen(AA)
print(meig$values)

## [1] 3.158981e+01 5.506463e+00 3.375284e+00 2.770820e+00 2.517158e+00
## [6] 2.388872e+00 2.317702e+00 2.277391e+00 2.256491e+00 8.582807e-06
cat("Vector for smallest ev:")

## Vector for smallest ev:
vecmin<-meig$vectors[,n]
print(vecmin)

## [1] 0.866015080 0.433009398 0.216509345 0.108264428 0.054151958 0.027115582
## [7] 0.013637056 0.006977088 0.003805678 0.002537115
cat("RQ for minimal eigenvector=",raynum(vecmin,AA))

## RQ for minimal eigenvector= 8.582807e-06
x<-proj(x) # Need to have parameters feasible
tmax<-system.time(asprqmax<-spg(x, fn=raynum, project=proj, A=-AA,

```

```

control=list(trace=TRUE, triter=1,maxit=1000)))[[3]]

## iter: 0 f-value: -20.09423 pgrad: 0.8771025
## iter: 1 f-value: -31.46235 pgrad: 0.03087639
## iter: 2 f-value: -31.55143 pgrad: 0.01613852
## iter: 3 f-value: -31.58762 pgrad: 0.00362995
## iter: 4 f-value: -31.58967 pgrad: 0.0009427463
## iter: 5 f-value: -31.5898 pgrad: 0.0002548985
## iter: 6 f-value: -31.58981 pgrad: 7.104836e-05
## iter: 7 f-value: -31.58981 pgrad: 2.018526e-05
## iter: 8 f-value: -31.58981 pgrad: 5.803606e-06
asprqmax

## $par
## [1] 0.085465700 -0.002889881 -0.091145405 -0.176319834 -0.255536076
## [6] -0.326113850 -0.385669103 -0.432189677 -0.464103537 -0.480334777
##
## $value
## [1] -31.58981
##
## $gradient
## [1] 5.803606e-06
##
## $fn.reduction
## [1] 11.49558
##
## $iter
## [1] 8
##
## $feval
## [1] 9
##
## $convergence
## [1] 0
##
## $message
## [1] "Successful convergence"
cat("Maximal eigenvalue is calculated as ", -asprqmax$value,"\n")

## Maximal eigenvalue is calculated as 31.58981
cat("Compare eigen:",meig$values[1]," difference=", -asprqmax$value-meig$values[1], "\n")

## Compare eigen: 31.58981 difference= -5.409149e-09
xy <- rep(1/sqrt(n), n)
tmax2<-system.time(asprqmax2<-spg(xy, fn=raynum, project=proj, A=-AA,
control=list(trace=TRUE, triter=1,maxit=1000)))[[3]]

## iter: 0 f-value: -20.5 pgrad: 0.7776636
## iter: 1 f-value: -31.51629 pgrad: 0.02149796
## iter: 2 f-value: -31.56598 pgrad: 0.01227049
## iter: 3 f-value: -31.58786 pgrad: 0.003475818
## iter: 4 f-value: -31.58965 pgrad: 0.0009871069
## iter: 5 f-value: -31.5898 pgrad: 0.0002837933

```

```

## iter: 6 f-value: -31.58981 pgrad: 8.226663e-05
## iter: 7 f-value: -31.58981 pgrad: 2.397799e-05
## iter: 8 f-value: -31.58981 pgrad: 7.015653e-06

cat("maximal eigensolution: Value=",asprqmax2$value,"in time ",tmax2,"\n")

## maximal eigensolution: Value= 31.58981 in time 0.001

print(asprqmax2$par)

## [1] 0.085466806 -0.002889806 -0.091147788 -0.176323486 -0.255538009
## [6] -0.326114501 -0.385668490 -0.432188271 -0.464102668 -0.480333914

tmin<-system.time(asprqmin<-spg(x, fn=raynum, project=proj, A=AA,
  control=list(trace=TRUE, triter=1,maxit=1000)))[[3]]

## iter: 0 f-value: 20.09423 pgrad: 0.8738632
## iter: 1 f-value: 0.1610048 pgrad: 0.6373408
## iter: 2 f-value: 1.253068 pgrad: 1.072525
## iter: 3 f-value: 1.44168 pgrad: 0.7840171
## iter: 4 f-value: 1.016774 pgrad: 0.6186038
## iter: 5 f-value: 3.940719 pgrad: 0.7924
## iter: 6 f-value: 0.03187537 pgrad: 0.4797277
## iter: 7 f-value: 0.02150955 pgrad: 0.3810032
## iter: 8 f-value: 0.01218731 pgrad: 0.3041359
## iter: 9 f-value: 0.001266454 pgrad: 0.07078474
## iter: 10 f-value: 0.000528509 pgrad: 0.08489315
## iter: 11 f-value: 0.00165654 pgrad: 0.1904217
## iter: 12 f-value: 0.0001417218 pgrad: 0.02495577
## iter: 13 f-value: 0.0001184352 pgrad: 0.0224623
## iter: 14 f-value: 8.854465e-05 pgrad: 0.01922957
## iter: 15 f-value: 9.035708e-06 pgrad: 0.001780281
## iter: 16 f-value: 9.541349e-06 pgrad: 0.005057676
## iter: 17 f-value: 1.10221e-05 pgrad: 0.008389123
## iter: 18 f-value: 8.598903e-06 pgrad: 0.0003246327
## iter: 19 f-value: 8.595024e-06 pgrad: 0.0002693203
## iter: 20 f-value: 8.584687e-06 pgrad: 8.147665e-05
## iter: 21 f-value: 8.582937e-06 pgrad: 2.117746e-05
## iter: 22 f-value: 8.582811e-06 pgrad: 6.798477e-06

cat("minimal eigensolution: Value=",asprqmin$value,"in time ",tmin,"\n")

## minimal eigensolution: Value= 8.582811e-06 in time 0.004

# print(asprqmin$par)
# Compare
cat("Diff from value from eigen():", (asprqmin$value-meig$values[n]))

## Diff from value from eigen(): 4.538548e-12

# cat("Vector difference:"); asprqmin$par - meig$vectors[,n]
# cat("Max abs relative difference=",max(abs((asprqmin$value-meig$values[n])/(abs(meig$values)+1e-18))))

```

If we ignore the constraint, and simply perform the optimization, we can sometimes get satisfactory solutions, though comparisons require that we normalize the parameters post-optimization. We can check if the scale of the eigenvectors is becoming large by computing the norm of the final parameter vector. In tests on the Moler matrix up to dimension 100, none grew to a worrying size.

For comparison, we also ran a specialized Geradin routine as implemented in R by one of us (JN). This gave

equivalent answers, albeit more efficiently. For those interested, the Geradin routine is available as referenced in Nash (2012).

## The extended Rosenbrock function on the unit ball

The `adagio` package (Borchers (2022)) gives an extended version of the Rosenbrock banana-shaped valley problem. This becomes a sumscale problem if we constrain the parameters to be on the unit ball, that is, where the sum of squares of the parameters is 1.

```
library(alabama)
library(optimx)
library(nloptr)

##
## Attaching package: 'nloptr'

## The following object is masked from 'package:alabama':
##
##      auglag

library(BB)
#####
# Minimizing a function on the unit ball: Min f(x), s.t. ||x|| = 1
#
rosbkext.f <- function(x){
  p <- x
  n <- length(p)
  sum (100*(p[1:(n-1)]^2 - p[2:n])^2 + (p[1:(n-1)] - 1)^2)
}

heq <- function(x){
  1 - sum(x*x)
}

transform <- function(x){
  # transforms x into z such that ||z|| = 1
  p <- length(x)
  z <- rep(NA, p)
  z[1] <- cos(x[1])
  z[p] <- prod(sin(x[-p]))
  if (p > 2) z[2:(p-1)] <- cumprod(sin(x[1:(p-2)]))*cos(x[2:(p-1)])
  return(z)
}

rosbkext.t <- function(x){
  n <- length(x)
  p <- transform(x)
  sum (100*(p[1:(n-1)]^2 - p[2:n])^2 + (p[1:(n-1)] - 1)^2)
}

ProjSphere <- function(x){
  x/sqrt(sum(x*x))
}

ProjSpheresgn <- function(x){
  sign(x[1])*x/sqrt(sum(x*x))
}
```

```

}

n <- 10
set.seed(1234)
p0 <- runif(n, 0, 3)

# Unconstrained optimization with parameter transformation to satisfy unit-length constraint
#
ans <- optim(par=p0, fn=rosbkext.t, method="BFGS", control=list(maxit=1000))
proptimr(ans)

## Result ans ( -> ) calc. min. = 9.418838 at
## -29.00744    156.519    109.596    -59.93688    -22.61133    4.065798    0.7760579    -0.0073
## After 74 fn evals, and 45 gr evals and NA hessian evals
## Termination code is 0 :
##
## -----
system.time(ans2 <- alabama::auglag(p0, rosbkext.f, heq=heq, control.outer = list(trace=FALSE, kktchk=F

## user system elapsed
## 0.028 0.000 0.027

proptimr(ans2)

## Result ans2 ( -> ) calc. min. = 6.426809 at
## 0.7478999    0.5648426    0.327218    0.1165006    0.02342218    0.01045982    0.01001565
## After 694 fn evals, and 117 gr evals and NA hessian evals
## Termination code is 0 :
## Gradient: [1] 4.816521e-02 4.430184e-02 1.899406e-02 8.673540e-03 2.029531e-03
## [6] 2.142209e-03 7.416204e-04 8.101936e-04 8.108949e-04 1.135815e-05
##
## -----
system.time(ans3 <- nloptr::slsqp(p0, rosbkext.f, heq=heq))

## user system elapsed
## 0.026 0.017 0.015

proptimr(ans3)

## Result ans3 ( -> ) calc. min. = 6.426807 at
## 0.7478806    0.5648321    0.3272499    0.116557    0.02348033    0.01050433    0.01006523
## After fn evals, and gr evals and NA hessian evals
## Termination code is 4 : NLOPT_XTOL_REACHED: Optimization stopped because xtol_rel or xtol_abs (above
##
## -----
system.time(ans4 <- spg(p0, rosbkext.f, project=ProjSphere))

## iter: 0 f-value: 5363.139 pgrad: 0.3108113
## iter: 10 f-value: 6.761595 pgrad: 1.170664
## iter: 20 f-value: 6.426807 pgrad: 0.0008499551

## user system elapsed
## 0.004 0.008 0.005

```

```
proptimr(ans4)
```

```
## Result  ans4 ( -> ) calc. min. = 6.426807 at
## 0.7478804    0.564832    0.3272502    0.1165575    0.02348068    0.01050432    0.01006517
## After  fn evals, and  gr evals and NA hessian evals
## Termination code is 0 : Successful convergence
## Gradient:[1] 2.012061e-05
##
## -----
```

```
system.time(ans4a <- spg(p0, rosbkext.f, project=ProjSpheresgn))
```

```
## iter: 0 f-value: 5363.139 pgrad: 0.3108113
## iter: 10 f-value: 6.761595 pgrad: 1.170664
## iter: 20 f-value: 6.426807 pgrad: 0.0008499551

## user system elapsed
## 0.004 0.000 0.005
```

```
proptimr(ans4a)
```

```
## Result  ans4a ( -> ) calc. min. = 6.426807 at
## 0.7478804    0.564832    0.3272502    0.1165575    0.02348068    0.01050432    0.01006517
## After  fn evals, and  gr evals and NA hessian evals
## Termination code is 0 : Successful convergence
## Gradient:[1] 2.012061e-05
##
## -----
```

```
c(ans$value, ans2$value, ans3$value, ans4$value, ans4a$value)
```

```
## [1] 9.418838 6.426809 6.426807 6.426807 6.426807
```

```
sevmeth <- c("ncg", "nvm", "BFGS", "L-BFGS-B", "tnewt", "ucminf", "spg")
```

```
several <- opm(p0, fn=rosbkext.t, gr="grcentral", method=sevmeth, control=list(trace=0))
```

```
## Warning in BB::spg(par = spar, fn = efn, gr = egr, lower = slower, upper =
## supper, : Unsuccessful convergence.
```

```
sumrbk<-summary(several, order=value, par.select=1:5)
```

```
print(sumrbk)
```

##		p1 s1	p2 s2	p3 s3	p4 s4	p5 s5			
##	nvm	-7.0091180	-53.9602431	3.4947023	0.2599813	5.5715970			
##	ucminf	0.7259327	-0.5531680	3.4947023	0.2599813	0.7115884			
##	ncg	5.5572526	-22.5443166	3.4947023	0.2599812	0.7115869			
##	BFGS	-7.0090845	-53.9600862	3.4936038	0.2464820	5.6408459			
##	L-BFGS-B	0.7258983	-0.5530069	2.7895973	2.8951592	3.7844436			
##	spg	0.7258107	0.5526231	-0.3496154	2.9288467	3.5030706			
##	tnewt	-8.6916438	-52.8463130	0.3607824	0.2591798	0.6887679			
##		value	fevals	gevals	hevals	conv	kkt1	kkt2	xtime
##	nvm	6.426807	115	81	0	0	TRUE	FALSE	0.024
##	ucminf	6.426807	62	62	0	0	TRUE	FALSE	0.009
##	ncg	6.426807	921	354	0	0	TRUE	FALSE	0.063
##	BFGS	6.436763	78	41	0	0	TRUE	FALSE	0.007
##	L-BFGS-B	6.436766	352	352	0	0	TRUE	FALSE	0.073
##	spg	6.456637	1900	1503	0	1	FALSE	FALSE	0.349
##	tnewt	9.408793	292	291	0	0	TRUE	FALSE	0.054

```

stp <- function(xx, fn){
  # standardize parameters in an opm output data-frame and check fn
  nr <- dim(xx)[1]
  npar <- which(colnames(xx)=="value") - 1
  newxx <- xx[, 1:(npar+1)]
  rownames(newxx)<-rownames(xx)
  colnames(newxx)<-colnames(xx)[1:(npar+1)]
  for (ii in 1:nr){
    meth <- rownames(xx)[ii]
    upar <- coef(xx[ii, ])
    tpar <- transform(upar) # spherical transform
    fval <- fn(upar)
    newxx[ii, 1:npar] <- tpar
    newxx[ii, npar+1] <- fval
    cat(meth, " fval=", fval, " at "); print(tpar[1:5])
  }
  newxx
}
tt<-stp(several, rosbkext.t)

## ncg fval= 6.426807 at [1] 0.74788061 0.56483208 0.32724994 0.11655698 0.02348039
## nvm fval= 6.426807 at [1] 0.74788061 0.56483208 0.32724994 0.11655697 0.02348037
## BFGS fval= 6.436763 at [1] 0.74790285 0.56486547 0.32728669 0.11658214 0.02348588
## L-BFGS-B fval= 6.436766 at [1] 0.74790344 0.56486635 0.32728606 0.11657767 0.02347118
## tnewt fval= 9.408793 at [1] -0.74308067 0.56671363 0.33299006 0.12143978 0.02485856
## ucminf fval= 6.426807 at [1] 0.74788061 0.56483208 0.32724994 0.11655697 0.02348037
## spg fval= 6.456637 at [1] 0.74796158 0.56494437 0.32733520 0.11665340 0.02357039

print(tt)

##          p1          p2          p3          p4          p5          p6
## ncg      0.7478806 0.5648321 0.3272499 0.1165570 0.02348039 1.050437e-02
## nvm      0.7478806 0.5648321 0.3272499 0.1165570 0.02348037 1.050432e-02
## BFGS     0.7479028 0.5648655 0.3272867 0.1165821 0.02348588 1.050428e-02
## L-BFGS-B 0.7479034 0.5648663 0.3272861 0.1165777 0.02347118 1.053857e-02
## tnewt    -0.7430807 0.5667136 0.3329901 0.1214398 0.02485856 1.066550e-02
## ucminf    0.7478806 0.5648321 0.3272499 0.1165570 0.02348037 1.050432e-02
## spg      0.7479616 0.5649444 0.3273352 0.1166534 0.02357039 -4.991341e-08
##          p7          p8          p9          p10         value
## ncg      0.010065153 0.010052391 9.860523e-03 9.605723e-05 6.426807
## nvm      0.010065226 0.010052399 9.860579e-03 9.585220e-05 6.426807
## BFGS     0.010041703 0.009878878 1.889037e-07 -3.083038e-05 6.436763
## L-BFGS-B 0.010155391 0.009738523 6.684097e-07 -9.084156e-05 6.436766
## tnewt    0.010158957 0.010144358 9.947400e-03 9.842338e-05 9.408793
## ucminf    0.010065225 0.010052400 9.860580e-03 9.585004e-05 6.426807
## spg      0.008808479 0.001056753 -3.395910e-04 7.733199e-04 6.456637
#####

```

Here we note that the `spg()` method that works well with a suitable projection, does much less well on the unconstrained minimization of the transformed objective `rosbkext.t()`.

## The R-help example

As a final example, let us use our present techniques to solve the problem posed by Lanh Tran on R-help. We will use only a method that scales the parameters directly inside the objective function and not bother



with gradients for this small problem.

```
ssums<-function(x){
  n<-length(x)
  tt<-sum(x)
  ss<-1:n
  xx<-(x/tt)*(x/tt)
  sum(ss*xx)
}

cat("Try penalized sum\n")

## Try penalized sum
require(optimx)
st<-runif(3)
aos<-opm(st, ssums, gr="grcentral", method="MOST")

## Warning in opm(st, ssums, gr = "grcentral", method = "MOST"): 'snewtonm'
## removed from 'method' -- no hess()

# rescale the parameters
nsol<-dim(aos)[1]
for (i in 1:nsol){
  tpar<-aos[i,1:3]
  ntpar<-sum(tpar)
  tpar<-tpar/ntpar
  # cat("Method ",aos[i, "meth"]," gives fval =", ssums(tpar))
  aos[i, 1:3]<-tpar
}
summary(aos,order=value)[1:5,]

##           p1 s1           p2 s2           p3 s3      value fevals gevals hevals
## subplex 0.5454545 0.2727273 0.1818182 0.5454545 416      0      0
## nlminb 0.5454545 0.2727273 0.1818182 0.5454545 9      8      0
## ncg 0.5454545 0.2727273 0.1818182 0.5454545 15      7      0
## nvm 0.5454545 0.2727273 0.1818182 0.5454545 15      9      0
## Rvmmin 0.5454545 0.2727273 0.1818182 0.5454545 15      9      0
##      conv kkt1 kkt2 xtime
## subplex 0 TRUE FALSE 0.002
## nlminb 0 TRUE FALSE 0.000
## ncg 0 TRUE FALSE 0.000
## nvm 0 TRUE FALSE 0.001
## Rvmmin 0 TRUE FALSE 0.002

ssum<-function(x){
  n<-length(x)
  ss<-1:n
  xx<-x*x
  sum(ss*xx)
}

proj.simplex <- function(y) {
  # project an n-dim vector y to the simplex Dn
  # Dn = { x : x n-dim, 1 >= x >= 0, sum(x) = 1}
  # Ravi Varadhan, Johns Hopkins University
  # August 8, 2012
  n <- length(y)
```

```

sy <- sort(y, decreasing=TRUE)
csy <- cumsum(sy)
rho <- max(which(sy > (csy - 1)/(1:n)))
theta <- (csy[rho] - 1) / rho
return(pmax(0, y - theta))
}
as<-spg(st, ssum, project=proj.simplex)

## iter: 0 f-value: 1.314879 pgrad: 0.3785277
## iter: 10 f-value: 0.5454545 pgrad: 9.544587e-06
cat("Using project.simplex with spg: fmin=",as$value," at \n")

## Using project.simplex with spg: fmin= 0.5454545 at
print(as$par)

## [1] 0.5454512 0.2727297 0.1818191

```

Apart from the parameter rescaling, this is an entirely “doable” problem. Note that we can also solve the problem as a Quadratic Program using the `quadprog` package.

```

library(quadprog)
Dmat<-diag(c(1,2,3))
Amat<-matrix(c(1, 1, 1), ncol=1)
bvec<-c(1)
meq=1
dvec<-c(0, 0, 0)
ans<-solve.QP(Dmat, dvec, Amat, bvec, meq=0, factorized=FALSE)
ans

## $solution
## [1] 0.5454545 0.2727273 0.1818182
##
## $value
## [1] 0.2727273
##
## $unconstrained.solution
## [1] 0 0 0
##
## $iterations
## [1] 2 0
##
## $Lagrangian
## [1] 0.5454545
##
## $iact
## [1] 1

```

## Conclusion

Sumscale problems can present difficulties for optimization (or function minimization) codes. These difficulties are by no means insurmountable, but they do require some attention.

While specialized approaches are “best” for speed and correctness, a general user is more likely to benefit from a simpler approach of embedding the scaling in the objective function and rescaling the parameters before reporting them. We also note that the use of a projected gradient via `spg` from package `BB` works very

well, but the projection needs to be set up carefully, as with the use of the sign of the first element in dealing with the Rayleigh Quotient.

## References

- Borchers, Hans W. 2022. *Adagio: Discrete and Global Optimization Routines*. <https://CRAN.R-project.org/package=adagio>.
- Geradin, M. 1971. “The Computational Efficiency of a New Minimization Algorithm for Eigenvalue Analysis.” *J. Sound Vib.* 19: 319–31.
- Nash, John C. 1979. *Compact Numerical Methods for Computers: Linear Algebra and Function Minimisation*. Bristol: Adam Hilger.
- . 2012. “Timing Rayleigh Quotient minimization in R.”